

Chapter 3

ARM Assembly Language

3.1 ARM Assembly Language and programming environment

ARM assembly language

- ARM instructions are written in the form

Label Op-code operand1, operand2, operand3 //comment

- Consider the following example implementing a loop.

```
Test_5 ADD  r0, r1, r2      //TotalTime = Time + NewTime
        SUBS r7, #1 //Decrement loop counter
        BEQ  Test_5          //IF zero THEN goto Test_5
```

- The Label field is a user-defined label that can be used by other instructions to refer to that line.
- Any text following a double slash is regarded as a comment field and is ignored by the assembler.

ARM assembly language and CPUlator

- We use CPUlator Simulator as the programming environment in this course. More details about this simulator will be provided in Lab 1.
- Web-based simulator and no installation needed:
<https://cpulator.01xz.net/?sys=arm-de1soc>
- Consider this program that you will see in Lab 1:

```
.org    0x1000    // Start at memory location 1000
.text                               // Code section
.global _start
_start:
mov r0, #9          // Store decimal 9 in register r0
mov r1, #0xE        // Store hex E (decimal 14) in register r1
add r2, r1, r0      // Add the contents of r0 and r1 and put result in r2

// End program
_stop:
b _stop

.end
```

ARM assembly language and CPUlator

- Two types of statements:
 1. Assembler directives
 - Used to control the assembler
 - Do not generate machine code
 - Highlighted in red.
 - All directives in CPUlator starts with a period.
 2. Executable instructions
 - Do generate machine code to perform various operations
 - *Pseudoinstruction*: instruction available to the programmer but not part of the processor's instruction set. Shorthand expression that an assembler converts to appropriate machine code (e.g., `ldr r1, =0xFF896788` – discuss in more detail after covering the fundamentals).

Directives used in the sample program

- `.org`
 - Format: `.org memory_address`
 - tells the compiler where in memory to place the program.
- `.text`
 - identifies the code section of the program.
- `.global`
 - Format: `.global label`
 - identifies globally visible labels.
- `.end`
 - identifies the end of the source code and data.

Instructions used in the sample program

mov (Move)

- Puts a value into a register.
- Two formats:
 - `mov r0, #9` puts the (immediate) decimal value 9 into register r0.
 - `mov r1, r0` copies the value in register r0 into register r1.

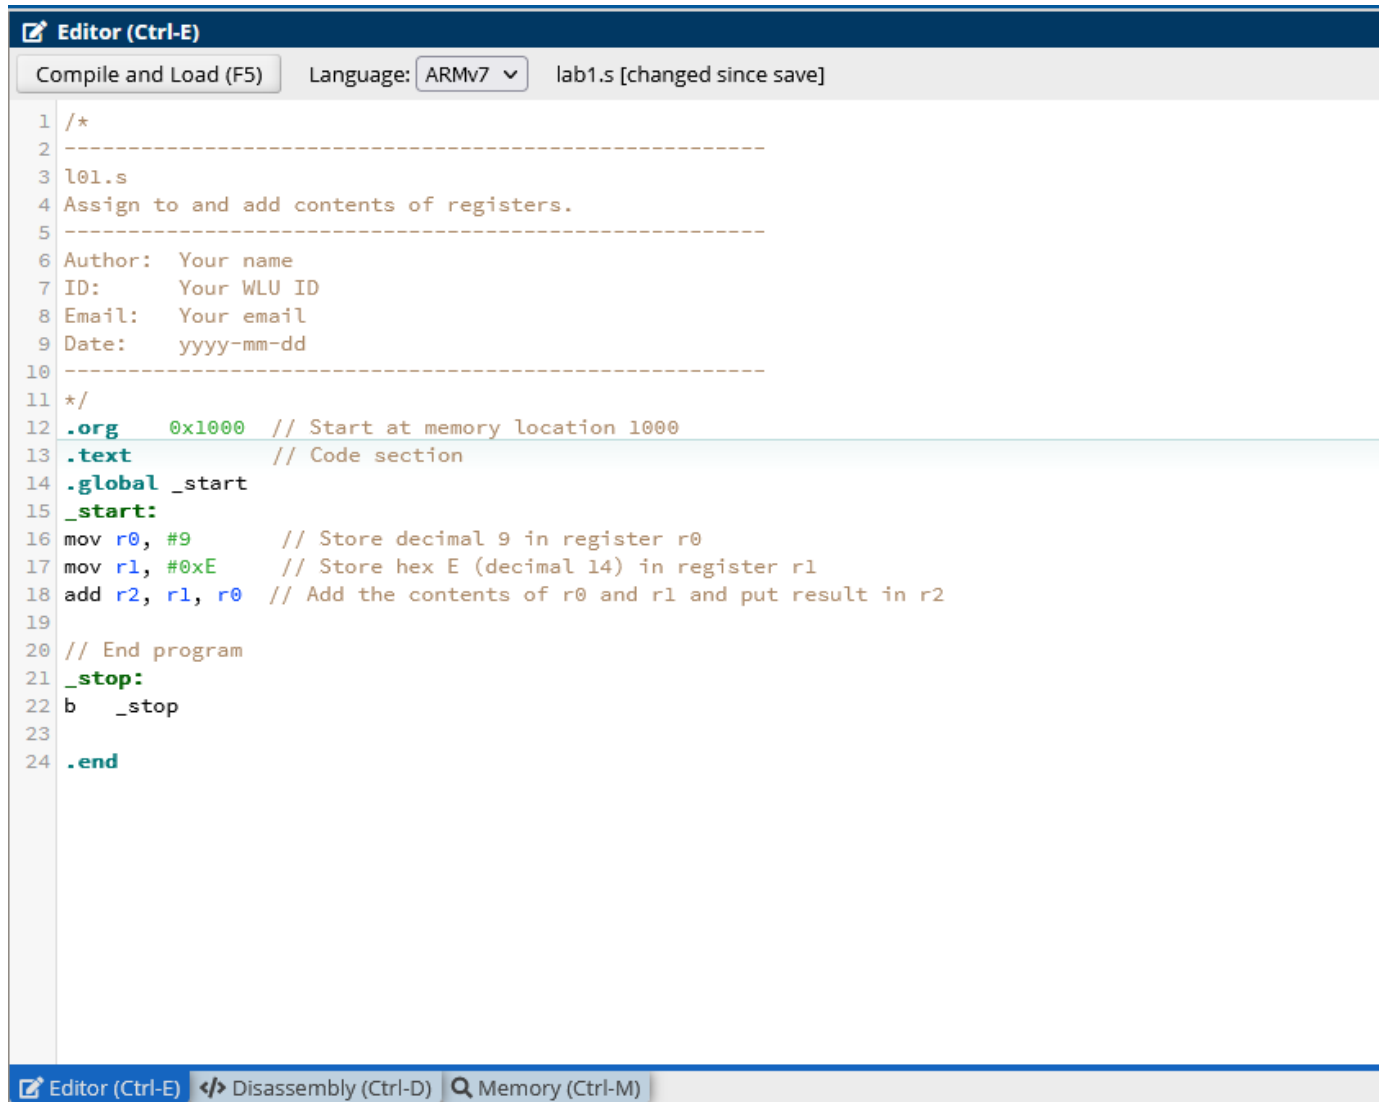
add (Add)

- adds two values and places the results into a register.
- Two formats:
 - `add r2, r1, r0` adds the contents of r0 and r1 and puts the result in register r2.
 - `add r3, r1, #7` adds the contents of r1 and the decimal value 7 and puts the result in register r3.

b (Branch)

- Branches execution to a memory location.
- The instruction:
 `_stop:`
 b `_stop`
terminates a program by branching to its own location.

CPUlator interface



The screenshot shows the CPUlator Editor window. The title bar is "Editor (Ctrl-E)". Below it is a toolbar with "Compile and Load (F5)", "Language: ARMv7", and "lab1.s [changed since save]". The main area contains assembly code for an ARMv7 program. The code is as follows:

```
1  /*
2  -----
3  lab1.s
4  Assign to and add contents of registers.
5  -----
6  Author:  Your name
7  ID:      Your WLU ID
8  Email:   Your email
9  Date:    yyyy-mm-dd
10 -----
11 */
12 .org    0x1000 // Start at memory location 1000
13 .text    // Code section
14 .global _start
15 _start:
16 mov r0, #9    // Store decimal 9 in register r0
17 mov r1, #0xE  // Store hex E (decimal 14) in register r1
18 add r2, r1, r0 // Add the contents of r0 and r1 and put result in r2
19
20 // End program
21 _stop:
22 b _stop
23
24 .end
```

At the bottom of the window is a navigation bar with "Editor (Ctrl-E)", "Disassembly (Ctrl-D)", and "Memory (Ctrl-M)".

CPUlator interface

Disassembly (Ctrl-D)		
Go to address, label, or register: <input type="text" value="00001000"/> Refresh		
Address	Opcode	Disassembly
00000ff0	00000000	andeq r0, r0, r0
00000ff4	00000000	andeq r0, r0, r0
00000ff8	00000000	andeq r0, r0, r0
00000ffc	00000000	andeq r0, r0, r0
11 */		
12 .org 0x1000 // Start at memory location 1000		
13 .text // Code section		
14 .global _start		
15 _start:		
16 mov r0, #9 // Store decimal 9 in register r0		
00001000	e3a00009	_start: mov r0, #9 ; 0x9
00001004	e3a0100e	17 mov r1, #0xE // Store hex E (decimal 14) in register r1 mov r1, #14 ; 0xe
00001008	e0812000	18 add r2, r1, r0 // Add the contents of r0 and r1 and put result in r2 add r2, r1, r0
20 // End program		
21 _stop:		
22 b _stop		
0000100c	eaffffff	_stop: b 0x100c (0x100c: _stop) _end:
00001010	aaaaaaaa	bge 0xfeaabac0
00001014	aaaaaaaa	bge 0xfeaabac4
00001018	aaaaaaaa	bge 0xfeaabac8
0000101c	aaaaaaaa	bge 0xfeaabacc
00001020	aaaaaaaa	bge 0xfeaabad0
00001024	aaaaaaaa	bge 0xfeaabad4
00001028	aaaaaaaa	bge 0xfeaabad8
0000102c	aaaaaaaa	bge 0xfeaabadc
00001030	aaaaaaaa	bge 0xfeaabae0
00001034	aaaaaaaa	bge 0xfeaabae4
00001038	aaaaaaaa	bge 0xfeaabae8

Address

- The address of the instruction.
- The program starts at address 0x1000 as specified by the .org instruction.

Opcode

- 32-bit instruction encoded in binary format
- Specify the operation (e.g., ADD) and the operands.

CPUlator interface

- You will learn the use of debugging tools CPUlator in Lab 2 and 3.
- The following statement you encounter in Lab 2 is a pseudoinstruction:
 - `ldr r0, =0x12345678`
 - This pseudoinstruction puts a number into a register
 - You will learn more why we need this statement and how this pseudoinstruction works later.

3.2 ARM Data-processing instructions

ARM Data Processing Instructions

- Consist of :
 - Arithmetic: **ADD** **ADC** **SUB** **SBC** **RSB** **RSC**
 - Logical: **AND** **ORR** **EOR** **BIC**
 - Comparisons: **CMP** **CMN** **TST** **TEQ**
 - Data movement: **MOV** **MVN**
- These instructions only work on registers, NOT memory.
- Syntax:
<Operation>{<cond>}{S} Rd, Rn, Operand2
 - Comparisons set flags only - they do not specify Rd
 - Data movement does not specify Rn
- Second operand is sent to the ALU via *barrel shifter*.

Updating ARM conditional codes

- ARM does not automatically update its status flags after an operation
- *Update-on-demand*: status flags are updated only the programmer puts an “S” suffix in the mnemonic, i.e.,
 - ADD r1, r2, r3 does not update status flags
 - ADD**S** r1, r2, r3 updates status flags
- This allows the programmer performs a test, and then carry out other instructions without changing the flags:

SUBS r1, r1, #1	; subtract 1 from r1 and set status bits
ADD r2, r2, #4	; increment r2 (don't update status bits)
BEQ Error	; if r1 is zero then deal with the problem


Addition

- Suppose you have a computer capable of adding 8-bit values but you want to add two 16-bit values.
- You can divide the 16-bit addition to two 8-bit additions, but the answer is not correct if there is a carry in the addition of the least significant byte:

Addition of two 8-bit number
without carry

$$\begin{array}{r} \text{34} \quad \text{32} \\ + \text{57} \quad \text{DF} \\ \hline \text{8B} \quad \text{11} \end{array}$$

Addition of two 8-bit number
with carry


$$\begin{array}{r} \text{34} \quad \text{32} \\ + \text{57} \quad \text{DF} \\ \hline \text{8C} \quad \text{11} \end{array}$$


Multi-Word Addition

- The following shows how a 64-bit addition is performed:

```
ADDS r4, r0, r2
```

```
ADC  r5, r1, r3
```



- Must use S after ADD, as the C flag in the status register (CPSR) must be updated.
- ADC: add with carry

Example


Analyze the following program which adds 0x35F62562FA to 0x21F412963B:

35F62562FA
+21F412963B

```
LDR    R0, =0xF62562FA    ; R0 = 0xF62562FA
LDR    R1, =0xF412963B    ; R1 = 0xF412963B
MOV     R2, #0x35          ; R2 = 0x35
MOV     R3, #0x21          ; R3 = 0x21
ADDS   R5, R1, R0          ; R5 = 0xF62562FA + 0xF412963B
                        ; now C = 1
ADC     R6, R2, R3          ; R6 = R2 + R3 + C
                        ;    = 0x35 + 21 + 1 = 0x57
```

Solution:

After the $R5 = R0 + R1$ the carry flag is one. Since $C = 1$, when ADC is executed, $R6 = R2 + R3 + C = 0x35 + 0x21 + 1 = 0x57$.



Microsoft Windows calculator support data size of up 64-bit (double word). Use it to verify the above calculations.

Subtraction

Subtraction:

- SUB $r1, r2, r3: [r1] \leftarrow [r2] - [r3]$
- SBC: subtract with borrow

Reverse subtraction:

- RSB $r1, r2, r3: [r1] \leftarrow [r3] - [r2]$
- As the literal must be in the second operand, RSB allows you to subtract the register content from a literal, e.g.,
 - RSB $r1, r2, \#10: [r1] \leftarrow 10 - [r2]$
- No negation operation in ARM. RSB also allows you to implement negation:
 - RSB $r1, r1, \#0: [r1] \leftarrow 0 - [r1]$

Examples: SUBS

Show the steps involved for the following cases:

a)

```
MOV    R2, #0x4F          ; R2 = 0x4F
MOV    R3, #0x39          ; R3 = 0x39
SUBS   R4, R2, R3          ; R4 = R2 - R3
```

b)

```
MOV    R2, #0x4F          ; R2 = 0x4F
SUBS   R4, R2, #0x05       ; R4 = R2 - 0x05
```

Solution:

a)

0x4F	0000004F	
- 0x39	+ FFFFFFFC7	2's complement of 0x39
16	1 00000016	

The flags would be set as follows: C = 1, Z = 0, N = 0 and V = 0

Note: if C = 1, No Borrow
if C = 0, Borrow

b)

0x4F	0000004F	
- 0x05	+ FFFFFFFB	2's complement of 0x05
0x4A	1 0000004A	

C = 1, Z = 0, N = 0 and V = 0

Example: Multi-Word Subtraction

Analyze the following program which subtracts 0x21F62562FA from 0x35F412963B:

35F412963B
- 21F62562FA

[illegible]

Solution:

After the $R5 = R1 - R0$ there is a borrow so the carry flag is cleared. Since $C = 0$, when SBC is executed, $R6 = R3 - R2 - 1 + C = 0x35 - 0x21 - 1 + 0 = 0x35 - 0x21 - 1 = 0x13$. **F412963B**

00000035
- 00000021

F412963B
- F62562FA

SBC R6,R3,R2 $\Rightarrow R6 = C - 1 + R3 - R2$
 $R6 = C - 1 + R3 + (2\text{'s complement of } R2)$

C = 0 so there is borrow

SUBS R5,R1,R0 =>
R5 = R1 + (2's complement of R0)

0	0	0	35
FF	FF	FF	DF
<hr/>			
0	0	0	13

+	F4	12	96	3B
	09	DA	9D	06
<hr/>				
0	FD	ED	33	41


Bitwise logical operations

Logical instruction ²⁵	Operation	Final value in r2
AND r2, r1, r0	11001010.00001111	00001010
OR r2, r1, r0	11001010+00001111	11001111
NOT r2, r1	$\overline{11001010}$	00110101
EOR r2, r1, r0	11001010 \oplus 00001111	11000101

- ARM does not have a NOT instruction of this form, but the NOT instruction can be implemented by one of the following two approaches:


1. mvn r2, r1: mvn copies the logical complement of r1 into r2
2. Apply EOR with the second operand equal to 0xFFFFFFFF. Note for each bit x , $x \text{ EOR } 1 = \bar{x}$
(i.e., 0 EOR 1 = 1
1 EOR 1 = 0)

Shift operations



(a) Logical shift


In a logical shift, a zero is shifted in and the bit shifted out is copied to the carry bit of the condition code register.



(b) Arithmetic shift

In an arithmetic shift, the number is either multiplied by 2 (ASL) or divided by 2 (ASR). The sign of a two's complement number is preserved.

The bit shifted out is copied into the carry bit.

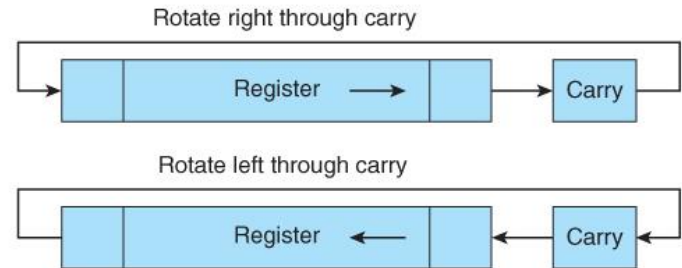


(c) Rotate

In a rotate operation, the bit shifted out is copied into the bit vacated at the other end (i.e., no bit is lost during a rotate). The bit shifted out is also copied into the carry bit.



The rotate through carry



© Cengage Learning 2014

- In eight bits, if the carry $C = 1$ and the word to be shifted is 01101110, a rotate left through carry would give
- 11011101 and carry = 0

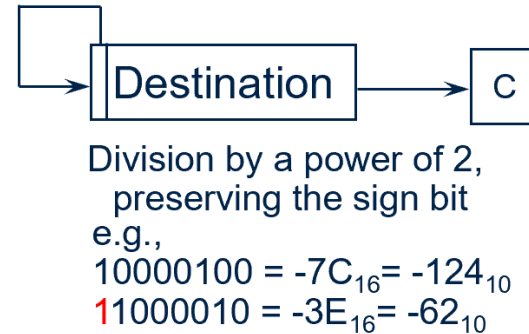
© Cengage Learning 2014

The Barrel Shifter in ARM

LSL : Logical Left Shift




ASR: Arithmetic Right Shift




LSR : Logical Shift Right



ROR: Rotate Right



RRX: Rotate Right Extended




Using the Barrel shifter: The second operand

- ARM combines shifting with other data processing operations, because the second operand can be shifted before it is used. Consider:

```
ADD r0, r1, r2, LSL #1
```

- A logical shift left is applied to the contents of r2 before they are added to the contents of r1. This operation is equivalent to $[r0] \leftarrow [r1] + [r2] \times 2$.
- Logical shift without other operations can be achieved by, e.g.:

```
MOV r0, r0, LSL #1
```



Static vs. Dynamic Shifts

Static shift

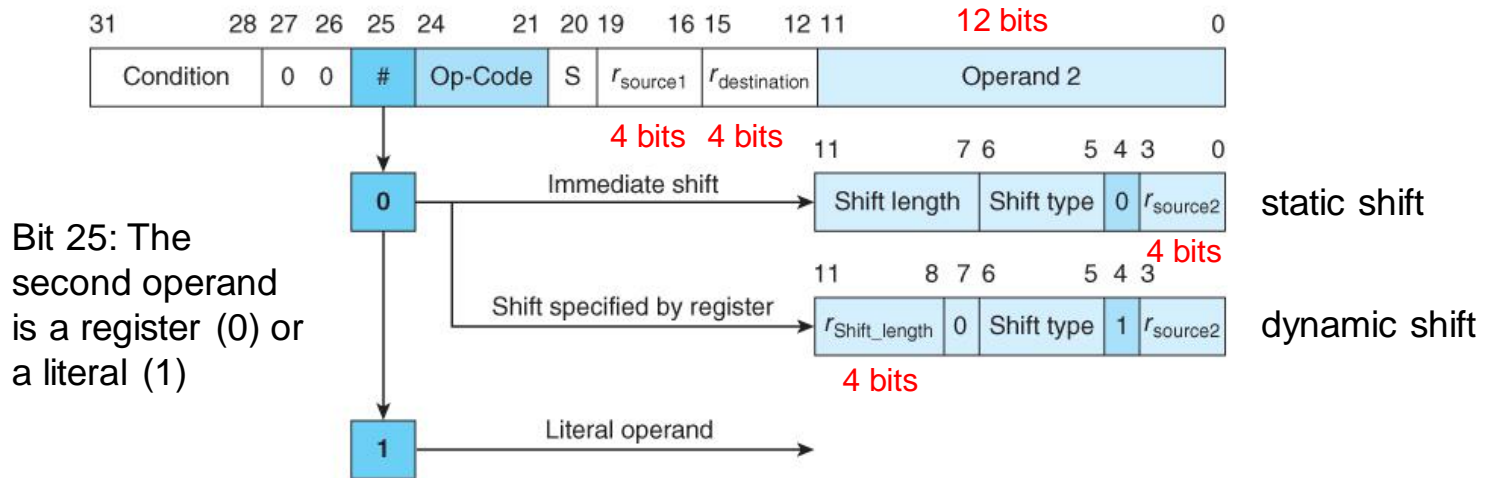
- The number of shifts is fixed in the code.
- e.g., `ADD r0, r1, r2, LSL #1` implements $[r0] \leftarrow [r1] + [r2] \times 2$

Dynamic shift

- Consider `ADD r0, r1, r2, LSL r3`, which implements $[r0] \leftarrow [r1] + [r2] \times 2^{[r3]}$
- This allows you to change the number of shifts at runtime.

Instruction Encoding

FIGURE 3.26 Encoding the ARM's data processing instructions




- 4 bits for each of the destination and first source operand.
- 12 bits for the second operand consisting of:
 - shift type
 - static shift: 5-bit unsigned integer in Bit 7 to 11 specifies the number of shifts. Range is 0 to 31.
 - dynamic shift: 4-bit address of the register in Bit 8 to 11 specifies the register containing the number of shifts.

Instruction Encoding


- Bit 21-24: Opcode

0000 = AND - $Rd := Op1 \text{ AND } Op2$
0001 = EOR - $Rd := Op1 \text{ EOR } Op2$
0010 = SUB - $Rd := Op1 - Op2$
0011 = RSB - $Rd := Op2 - Op1$
0100 = ADD - $Rd := Op1 + Op2$
0101 = ADC - $Rd := Op1 + Op2 + C$
0110 = SBC - $Rd := Op1 - Op2 + C - 1$
0111 = RSC - $Rd := Op2 - Op1 + C - 1$
1000 = TST - set condition codes on $Op1 \text{ AND } Op2$
1001 = TEQ - set condition codes on $Op1 \text{ EOR } Op2$
1010 = CMP - set condition codes on $Op1 - Op2$
1011 = CMN - set condition codes on $Op1 + Op2$
1100 = ORR - $Rd := Op1 \text{ OR } Op2$
1101 = MOV - $Rd := Op2$
1110 = BIC - $Rd := Op1 \text{ AND NOT } Op2$
1111 = MVN - $Rd := \text{NOT } Op2$

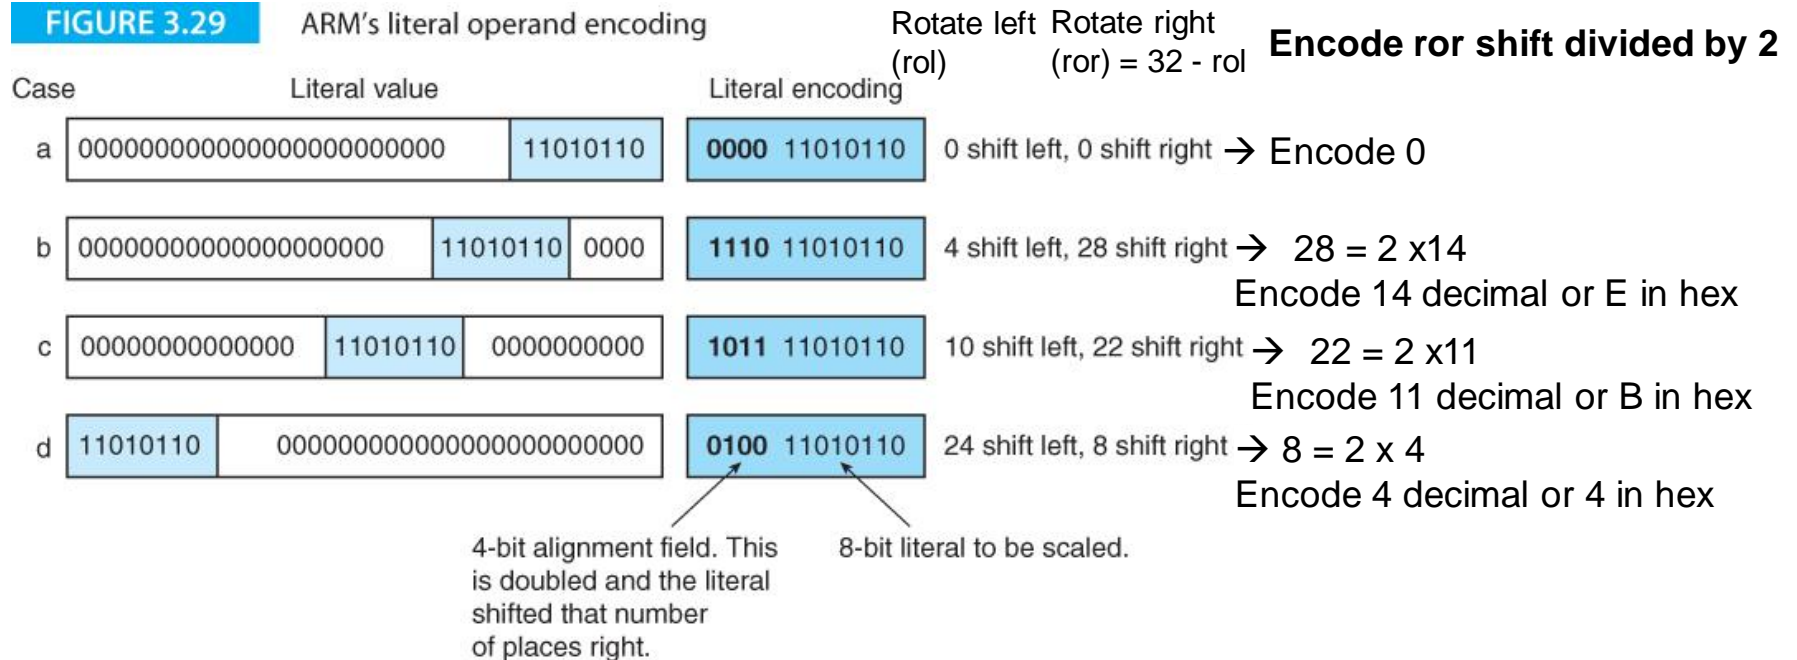
Literal Encoding



- No ARM instruction can contain a 32-bit immediate constant
- The data processing instruction format has 12 bits available for operand2
- The literal processed by the instruction is the 8-bit value stored in Bits 0-7, shifted by ROR by 2 times the value stored in Bits 8-11.



Example of literal encoding



Example of literal encoding

- The assembler converts immediate values to the rotate form:
 - `MOV r0,#4096;` uses `0x01`, left shift 12, ror 20, encode 10 or `A`
machine code: `E3A00A01`
 - `ADD r1,r2,#0xFF0000;` uses `0xFF`, left shift 16, ror 16, encode `8`
machine code: `E28218FF`
- Should be able to convert least significant 12 bit of machine code to the literal being encoded
- The MVN (move negative) instruction `MVN r0, #XX` can specify unshifted constant in the range `0xFFFFFFFF00` to `0xFFFFFFFFFF`. e.g.,
 - `MOV r0, #0xFFFFFFFF ;assembles to MVN r0,#0`
- Values that cannot be generated in this way will cause an error.

More examples on literal encoding

Determine the literal operand encoded by the following 32-bit machine code. Express your answer in hexadecimal representation.

(a) E3A00C67

(b) E2810E89

(c) E283260F

(d) E3A04807

More examples on literal encoding

Determine the literal operand encoded by the following 32-bit machine code. Express your answer in hexadecimal representation.

(a) E3A00C67

Encoded C \rightarrow $\text{ror } 12 \times 2 = 24 \rightarrow$ left shift $= 32 - 24 = 8$.

Literal = 67 shifted by 8 bits to the left = 6700

(b) E2810E89

Encoded E \rightarrow $\text{ror } 14 \times 2 = 28 \rightarrow$ left shift $= 32 - 28 = 4$.

Literal = 89 shifted by 4 bits to the left = 890

(c) E283260F

Encoded 6 \rightarrow $\text{ror } 6 \times 2 = 12 \rightarrow$ left shift $= 32 - 12 = 20$.

Literal = 0F shifted by 20 bits to the left = F00000

(d) E3A04807

Encoded 8 \rightarrow $\text{ror } 8 \times 2 = 16 \rightarrow$ left shift $= 32 - 16 = 16$.

Literal = 07 shifted by 16 bits to the left = 70000

Multiplication

- Take two m -bit operands \rightarrow output a $2m$ -bit result.
- Multiplication would not be correct using 2's complement values.
- Cannot use the same multiplication operation in signed and unsigned values. Focus on unsigned here.
- For more information:
https://bohr.wlu.ca/cp216/docs/QRC0001_UAL.pdf

Unsigned multiplications

Operation	Assembler	Action
Multiply and accumulate and subtract unsigned long unsigned accumulate long	<code>MUL{S} Rd, Rm, Rs</code> <code>MLA{S} Rd, Rm, Rs, Rn</code> <code>MLS Rd, Rm, Rs, Rn</code> <code>UMULL{S} RdLo, RdHi, Rm, Rs</code> <code>UMLAL{S} RdLo, RdHi, Rm, Rs</code>	$Rd := (Rm * Rs)[31:0]$ $Rd := (Rn + (Rm * Rs))[31:0]$ $Rd := (Rn - (Rm * Rs))[31:0]$ $RdHi, RdLo := \text{unsigned}(Rm * Rs)$ $RdHi, RdLo := \text{unsigned}(RdHi, RdLo + Rm * Rs)$

- Cannot use the same register to specify both `Rd` and `Rm`.
- `MUL`, `MLA` and `MLS` have the multiplication results truncated to the lower-order 32 bits.
- `UMULL` and `UMLAL` store the full 64-bit result into two registers `RdLo` and `RdHi`.

3.3 ARM Flow Control Instructions

Unconditional branch

- **Format:** `B target`, where `target` denotes the branch target address (the address of the next instruction to be executed).

- **Example:**

```
..    do this      Some code
..    then that    Some other code
      B    Next    Now skip past next instructions
..        ...the code being skipped past
..        ...the code being skipped past
Next ..    Target address for the branch
```

Conditional branch

```
IF (X == Y)
    THEN Y = Y + 1;
    ELSE Y = Y + 2
```

- A test is performed and one of two courses of action is carried out depending on the outcome. We can translate this as:

```
    CMP r1,r2      ;r1 contains y and r2 contains x: compare them
    BNE Plus2      ;if not equal, then branch to the else part
    ADD r1,r1,#1   ;if equal, fall through to here and + 1 to y
    B leave        ;now skip past the else part
Plus2 ADD r1,r1,#2 ;ELSE part add 2 to y
leave ...         ;continue from here
```

- The *conditional branch* instruction tests flag bits in the processor's CPSR, then takes the branch if the tested condition is true.
- Here `CMP` computes $[r1] - [r2]$ and update the CPSR, the conditional branch tests for the condition NE (not equal), corresponding to !Z, i.e., condition is true when $Z = 0$.
- `CMP` always updates the CPSR flags (i.e., no S suffix is needed)

Condition Codes

The code of possible conditions is tabulated here:

Code	Suffix	Description	Flags
0000	EQ	Equal / equals zero	Z
0001	NE	Not equal	!Z
0010	CS / HS	Carry set / unsigned higher or same	C
0011	CC / LO	Carry clear / unsigned lower	!C
0100	MI	Minus / negative	N
0101	PL	Plus / positive or zero	!N
0110	VS	Overflow	V
0111	VC	No overflow	!V
1000	HI	Unsigned higher	C and !Z
1001	LS	Unsigned lower or same	!C or Z
1010	GE	Signed greater than or equal	N == V
1011	LT	Signed less than	N != V
1100	GT	Signed greater than	!Z and (N == V)
1101	LE	Signed less than or equal	Z or (N != V)
1110	AL	Always (default)	any

Branching for signed and unsigned data

- Branch instruction can refer to signed and unsigned data
- Consider the example:
 - $[r1] = 0x90000075$
 - $[r2] = 0x50000075$
 - Which value is larger? The answer depends on whether **you interpret** the number as signed or unsigned numbers.
 - if interpreted as unsigned numbers, $[r1] > [r2]$.
 - if interpreted as signed numbers, $[r1]$ is a negative number and $[r2]$ is a positive number. Thus $[r2] > [r1]$.

Branching for signed and unsigned data

- The computer does not know whether you want to interpret the comparison as signed and unsigned.
- When executing `CMP r1, r2`, the computer would just perform $[r1] - [r2]$ and update CPSR.

90000075	90000075
- 50000075	<u>+AFFFFFF8B</u> (2's complement)
	1 40000000

CPSR: C = 1, N = 0, V = 1 (-ve + -ve = +ve), Z = 0

- If programmer decides to test $[r1] > [r2]$ and consider the operands as *unsigned numbers* → Use BHI (Unsigned higher)
 - BHI branches as the condition C and !Z are satisfied
- If programmer decides to test $[r1] > [r2]$ and consider the operands as *signed numbers* → Use BGT (Signed greater than)
 - BGT does not take the branch as the condition $N \neq V$

Test instruction: TST

Two test instructions that explicitly update CPSR:

`TST Rn, Operand2`

- updates CPSR on Rn AND Operand2
- Useful in testing whether an individual bit of a word is 1.
- E.g., Bit 5 of a lower-case ASCII character is set to 1. To test whether an ASCII character in r0 is lower-case:

```
TST r0, #2_0010000
```

```
BNE LowerCase // jump to code handling lower case
```

- If lower-case, then the result is non-zero (i.e., Z = 0) → Branch taken
- If capital, result is zero (i.e., Z = 1) → Branch not taken


Test instruction: TEQ

TEQ Rn, <Operand2>

- updates CPSR on Rn EOR Operand2
- Z-bit is high if [Rn] == [Operand2]
 - [Rn] == [Operand2]: All bits in the result are 0 (i.e., Z = 1)
 - [Rn] != [Operand2]: At least one bit in the result are 1 (i.e., Z = 0)
- Similar to CMP but TEQ does not update the overflow flag (no concept of V in logical operations).
- Only used to test equivalence, not less/larger than conditions.

Conditional Execution

- In ARM, each instruction is conditionally executed (i.e., an instruction is executed only if a condition after the conditional code list is satisfied)
- So far, we have only dealt with the default case *always*:
 - ADD r0, r1, r2 means ADDAL r0, r1, r2
- Bits 28-31 in Fig. 3.26 are used to encode the condition.



- e.g., if $Z = 1$ then $[r0] \leftarrow [r1] + [r2]$ can be implemented as: ADDEQ r0, r1, r2

Example 1

- This improves code density and performance by reducing the number of forward branch instructions.
- E.g., if (a == 0); b = c + d

```
CMP    r3, #0
BNE     skip
ADD     r0, r1, r2
skip
```

```
CMP    r3, #0
ADDEQ   r0, r1, r2
```

Example 2

```
if (a == b) e = e + 4;
if (a < b)   e = e + 7;
if (a > b)   e = e + 12;
```

Using the same register assignments as before, we can use conditional execution to implement this as

```
CMP    r0,r1      ;compare a == b
ADDEQ  r4,r4,#4    ;if a == b then e = e + 4
ADDLE  r4,r4,#7    ;if a < b then e = e + 7
ADDGT  r4,r4,#12   ;if a > b then e = e + 12
```

Once again, using conventional non-conditional execution, we would have to write something like the following to implement this algorithm.

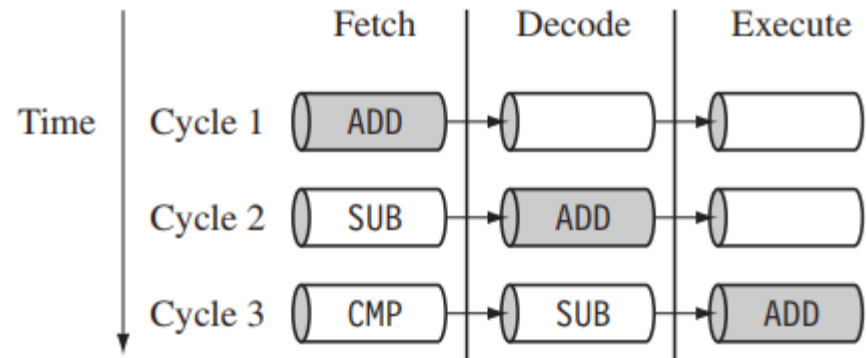
	CMP	r0,r1				;compare a == b
	BNE	Test1				;not equal try next test
	ADD	r4,r4,#4				;a == b so e = e+4
	B	ExitAll				;now leave
Test1	BLT	Test2				;if a < b then
	ADD	r4,r4,#12				;if we are here a > b so e = e + 12
	B	ExitAll				;now leave
Test2	ADD	r4,r4,#7				;if we are here a < b so e = e + 7
ExitAll						

a==b a<b a>b

Pipelining

- ARM7 has a three-stage pipeline:

- fetch
- decode
- execute



- When ADD is being executed (Cycle 3), the pc is pointing to CMP, which is two instructions (8 bytes) below ADD.

- pc = executed address + 8









- For our example, when ADD is executed:

- pc = 0x100C + 8 = 0x1014

Address	Instruction
100C	ADD
1010	SUB
pc → 1014	CMP

No Pipelining

30 30 30 30 30 30 30 30



With Pipelining









30

30

30

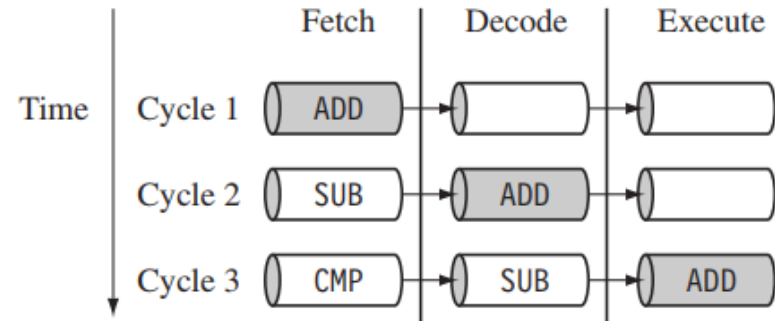
30

30



Pipelining

- When ADD is being executed (Cycle 3), the PC is pointing to CMP, which is



two instructions (8 bytes) below ADD.

- $PC = \text{executed address} + 8$
- This has an effect on address encoding, as described below.


	Address	Instruction
	100C	ADD
	1010	SUB
pc →	1014	CMP

Branch instruction encoding

- How does the machine code encode the destination address of the branching operation?

FIGURE 3.41


Encoding ARM's branch and branch-with-link instructions



Branch instruction encoding

FIGURE 3.41

Encoding ARM's branch and branch-with-link instructions

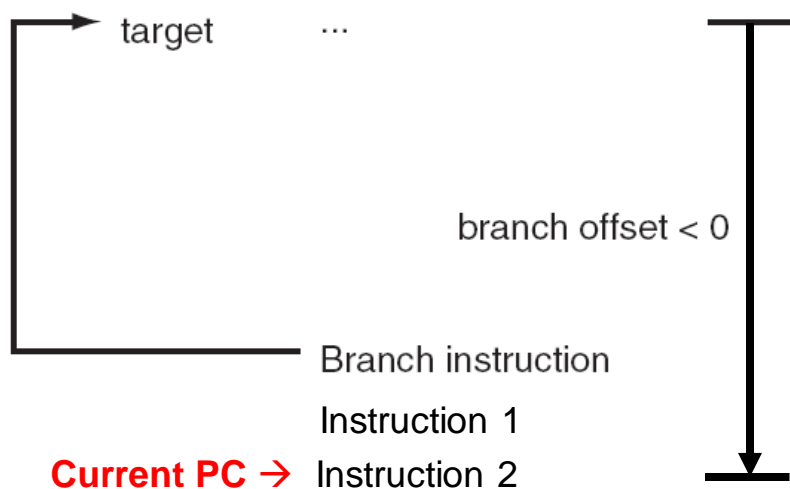


- The 24-bit signed offset stores the number of instructions of the destination instruction with respect to the current PC.
- Note each instruction has 4 bytes.
- Thus, the 24-bit offset is shifted left twice to convert the instruction offset to byte offset → 26-bit signed offset
- Range of offset: -32 Mbytes to +32 Mbytes or -8 M instructions to +8 M instructions

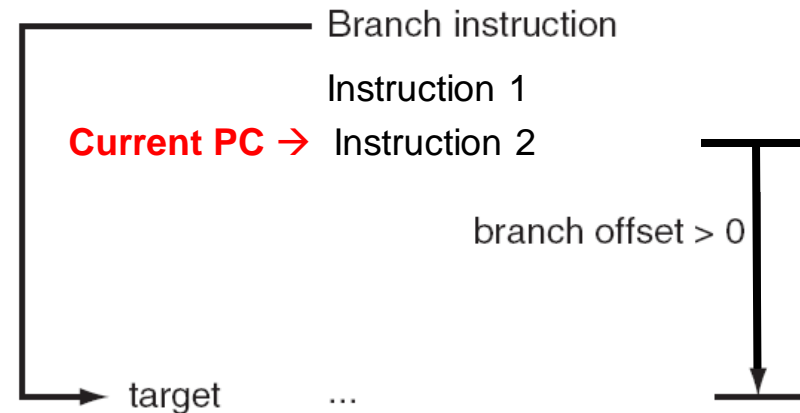
26-bit signed: -2^{25} to $2^{25}-1$ bytes

24-bit signed: -2^{23} to $2^{23}-1$ instructions $1\text{M} = 2^{20}$

Branch offset



(a) Backward branch (-ve offset)



(a) Forward branch (+ve offset)

- Note that due to pipelining, when a branch instruction is executed, the PC points to two instructions below the branching instruction.

Example 1

Address	Machine	_start:
00000000	e3a00000	23 mov r0, #0 ; 0x0
00000004	e3a02ffa	24 ldr r2, =1000 // Step 1 mov r2, #1000 ; 0x3e8
00000008	e2800009	25 AddMore: 26 add r0, r0, #9 // Step 2 AddMore: add r0, r0, #9 ; 0x9
0000000c	e2522001	27 subs r2, r2, #1 // Step 3 subs r2, r2, #1 ; 0x1
00000010	1affffff	28 bne AddMore bne 0x8 (0x8: AddMore)
00000014	e1a04000	29 mov r4, r0
00000018	eaffffff	31 Here: 32 b Here Here: b 0x18 (0x18: Here)

$$\begin{array}{r}
 000004 \\
 \text{FFFFFB} \\
 + \quad 1 \\
 \hline
 \text{FFFFFC}
 \end{array}$$

2's complement

$$\begin{aligned}
 &\text{Encoded address} \\
 &= \frac{\text{Dest} - \text{Current PC}}{4}
 \end{aligned}$$

- Dest = 0x08
- Current PC = 0x18
(two instructions beyond bne)

$$\begin{aligned}
 &\text{Encoded address} = \\
 &\frac{0x08 - 0x18}{4} = -4 \\
 &= 0xFFFFFC
 \end{aligned}$$

Example 2

Determine the encoded destination addresses in the covered machine code:

```

4 mov r0, #7
5 mov r1, #-8
6 mov r4, #0
7
8 // Non-conditional execution
9 cmp r0, r1
10 beq Label1
11 bgt Label2
12 add r4, r4, #1
13 b ExitAll
14
15 Label2:
16 add r4, r4, #4
17 b ExitAll
18
19 Label1:
20 add r4, r4, #2
21
22
23 ExitAll:
24 b ExitAll
25
26 .end

```

Address	Opcode	Disassembly
00000000	e3a00007	_start: 4 mov r0, #7 ; 0x7
00000004	e3e01007	5 mov r1, #-8 mvn r1, #7 ; 0x7
00000008	e3a04000	6 mov r4, #0 ; 0x0
0000000c	e1500001	8 // Non-conditional execution 9 cmp r0, r1
00000010	0a	10 beq Label1 beq 0x28 (0x28: Label1)
00000014	ca	11 bgt Label2 bgt 0x20 (0x20: Label2)
00000018	e2844001	12 add r4, r4, #1 ; 0x1
0000001c	ea	13 b ExitAll b 0x2c (0x2c: ExitAll)
00000020	e2844004	15 Label2: Label2: 16 add r4, r4, #4 ; 0x4
00000024	ea	17 b ExitAll b 0x2c (0x2c: ExitAll)
00000028	e2844002	19 Label1: Label1: 20 add r4, r4, #2 ; 0x2
0000002c	ea	23 ExitAll: 24 b ExitAll ExitAll: b 0x2c (0x2c: ExitAll) _end:

Example 2

Address	Opcode	Disassembly
00000000	e3a00007	_start: 4 mov r0, #7 ; 0x7
00000004	e3e01007	5 mov r1, #-8
00000008	e3a04000	6 mov r4, #0 ; 0x0
0000000c	e1500001	8 // Non-conditional execution 9 cmp r0, r1
00000010	0a000004	10 beq Label1 beq 0x28 (0x28: Label1)
00000014	ca000001	11 bgt Label2 bgt 0x20 (0x20: Label2)
00000018	e2844001	12 add r4, r4, #1 ; 0x1
0000001c	ea000002	13 b ExitAll b 0x2c (0x2c: ExitAll)
00000020	e2844004	15 Label2: Label2: 16 add r4, r4, #4 ; 0x4
00000024	ea000000	17 b ExitAll b 0x2c (0x2c: ExitAll)
00000028	e2844002	19 Label1: Label1: 20 add r4, r4, #2 ; 0x2
0000002c	ea000000	23 ExitAll: 24 b ExitAll ExitAll: b 0x2c (0x2c: ExitAll) _end:

Example 2

Address	Opcode	Disassembly
		_start:
00000000	e3a00007	4 mov r0, #7 ; 0x7
		5 mov r1, #-8
00000004	e3e01007	6 mvn r1, #7 ; 0x7
00000008	e3a04000	7 mov r4, #0 ; 0x0
		8 // Non-conditional execution
0000000c	e1500001	9 cmp r0, r1
		10 beq Label1
00000010	0a000004	11 beq 0x28 (0x28: Label1)
		12 bgt Label2
00000014	ca000001	13 bgt 0x20 (0x20: Label2)
00000018	e2844001	14 add r4, r4, #1 ; 0x1
		15 b ExitAll
0000001c	ea000002	16 b 0x2c (0x2c: ExitAll)
		17 Label2:
00000020	e2844004	18 add r4, r4, #4 ; 0x4
		19 b ExitAll
00000024	ea000000	20 b 0x2c (0x2c: ExitAll)
		21 Label1:
00000028	e2844002	22 add r4, r4, #2 ; 0x2
		23 ExitAll:
		24 b ExitAll
		ExitAll:
0000002c	ea000000	25 b 0x2c (0x2c: ExitAll)
		_end:

00000034 // add 8
from address to get
PC

0000002
(-2) FFFFFD
Encode: + FFFFFFFE

Address 0x10

Dest = 0x28, Current PC = 0x18

$$R = \frac{0x28 - 0x18}{4} = 4$$

Encode: 000004

Address 0x14

Dest = 0x20, Current PC = 0x1c

$$R = \frac{0x20 - 0x1c}{4} = 1$$

Encode: 000001

Address 0x1c

Dest = 0x2c, Current PC = 0x24

$$R = \frac{0x2c - 0x24}{4} = 2$$

Encode: 000002

Address 0x24

Dest = 0x2c, Current PC = 0x2c

$$R = \frac{0x2c - 0x2c}{4} = 0$$

Encode: 000000

Address 0x2c

Dest = 0x2c, Current PC = 0x34

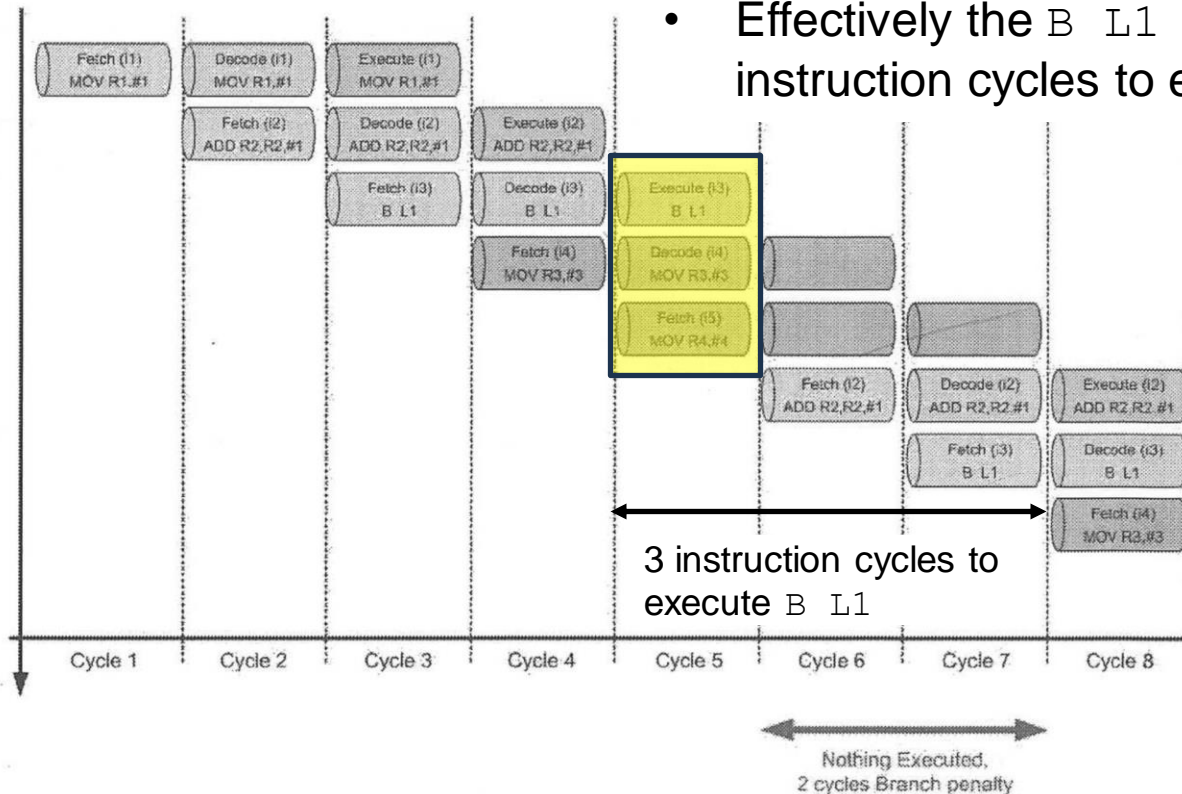
$$R = \frac{0x2c - 0x34}{4} = -2$$

55

Branch penalty

```
MOV R1, #1
L1:
ADD R2, R2, #1
B L1
MOV R3, #3
MOV R4, #4
```

- First instruction takes 3 cycles to pass through the stages of the pipeline
- Thereafter, instructions take one cycle to execute
- The B L1 instruction changes the sequential order of program execution.
- The instruction in the fetch and decode stages must be dumped (or flushed).
- Effectively the B L1 instruction takes 3 instruction cycles to execute.



Example

```
mov r0, #255
```

```
Again: subs r0, r0, #1 //1 cycle
```

```
bne Again           //1 cycle if not  
                    branch, 3 if branch
```

How many instruction cycles are generated by the Again loop?

	[r0]	Instruction Cycles
1 st repetition	255 → 254	4
2 nd repetition	254 → 253	4
⋮	⋮	⋮
255 th repetition	1 → 0	2 (bne branch not taken)
	Total	1018

254 x 4